

unit-2assign-2

November 18, 2024

```
[1]: import pandas as pd

# Load the dataset
data = pd.read_csv("NvWCPg2u12Tza1js.csv")

# Preview the dataset
print(data.head())

# Summary statistics
print(data.describe())

# Check for missing values
print(data.isnull().sum())
```

	Age	Education	Race	Hisp	MaritalStatus	Nodeg	\
0	45	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
1	21	Intermediate	NotBlack	NotHispanic	NotMarried	0	
2	38	HighSchool	NotBlack	NotHispanic	Married	0	
3	48	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
4	18	LessThanHighSchool	NotBlack	NotHispanic	Married	1	

	Earnings_1974	Earnings_1975	Earnings_1978
0	21516.670	25243.550	25564.670
1	3175.971	5852.565	13496.080
2	23039.020	25130.760	25564.670
3	24994.370	25243.550	25564.670
4	1669.295	10727.610	9860.869

	Age	Nodeg	Earnings_1974	Earnings_1975	Earnings_1978
count	15992.000000	15992.000000	15992.000000	15992.000000	15992.000000
mean	33.225238	0.295835	14016.800304	13650.803376	14846.659673
std	11.045216	0.456432	9569.795893	9270.403225	9647.391524
min	16.000000	0.000000	0.000000	0.000000	0.000000
25%	24.000000	0.000000	4403.452250	4398.823000	5669.298000
50%	31.000000	0.000000	15123.580000	14557.110000	16421.975000
75%	42.000000	1.000000	23584.180000	22923.737500	25564.670000
max	55.000000	1.000000	25862.320000	25243.550000	25564.670000
Age	0				
Education	0				

```

Race          0
Hisp          0
MaritalStatus 0
Nodeg         0
Earnings_1974 0
Earnings_1975 0
Earnings_1978 0
dtype: int64

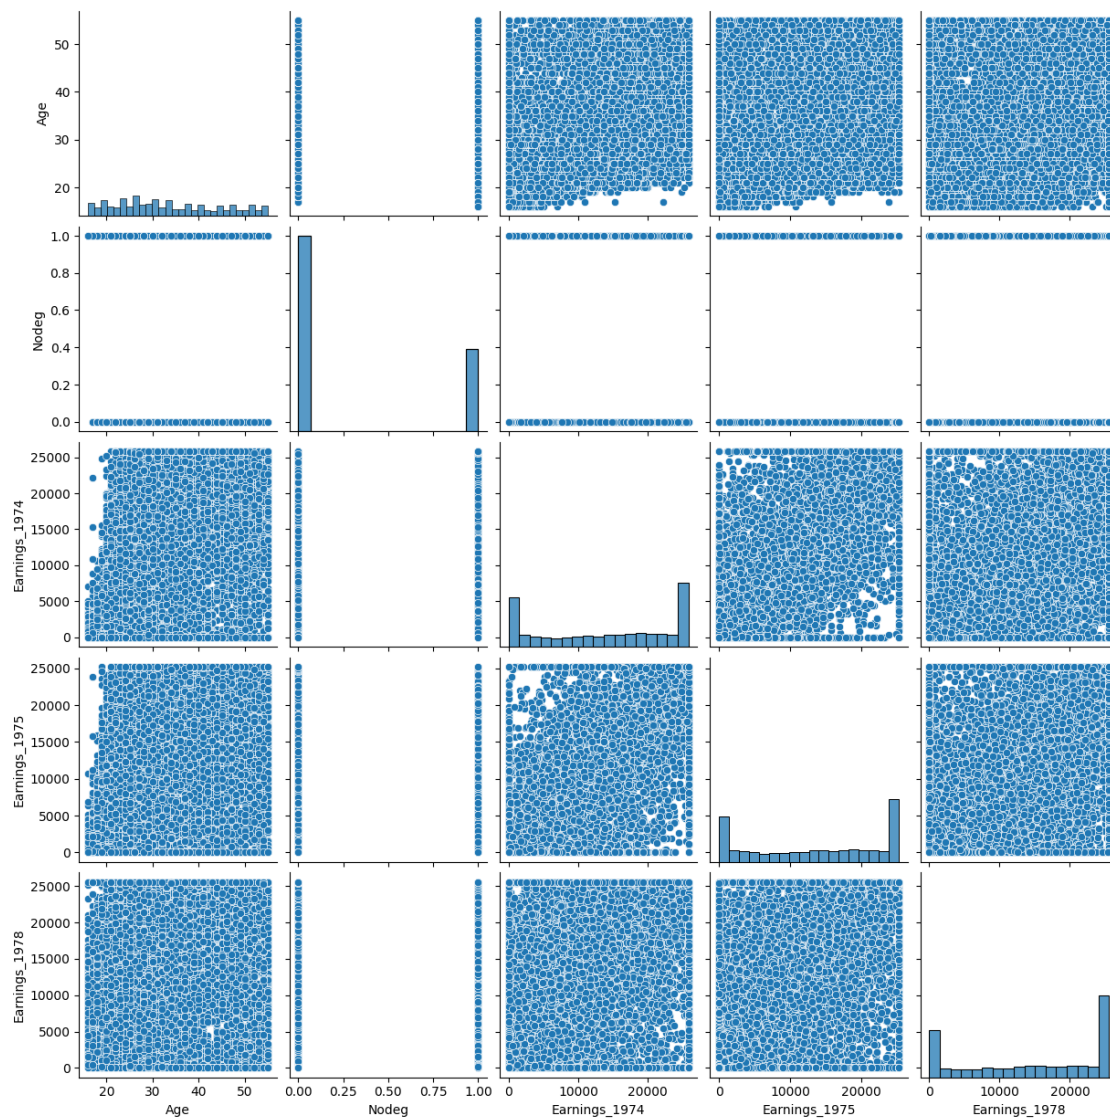
```

```

[2]: import seaborn as sns
import matplotlib.pyplot as plt

# Plot pairwise relationships
sns.pairplot(data)
plt.show()

```



```
[8]: X = X.apply(pd.to_numeric, errors='coerce') # Converts non-numeric values to
      ↪ NaN
      print(X.isnull().sum()) # Check for any NaN introduced
      X = X.fillna(0) # Replace NaN with 0
```

```
Age                                0
Nodeg                             0
Education_Intermediate             0
Education_LessThanHighSchool      0
Education_PostGraduate            0
Education_graduate                 0
Race_black                        0
Hisp_hispanic                     0
MaritalStatus_NotMarried          0
dtype: int64
```

```
[11]: from sklearn.model_selection import train_test_split

      # Prepare data
      # Replace 'Earnings_1978' with the desired target variable
      X = data.drop(['Earnings_1974', 'Earnings_1975', 'Earnings_1978'], axis=1)
      y = data['Earnings_1978'] # Use the appropriate column as the target variable

      # Split the dataset into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)

      # Display the shapes of the splits to verify
      print(f"X_train shape: {X_train.shape}")
      print(f"X_test shape: {X_test.shape}")
      print(f"y_train shape: {y_train.shape}")
      print(f"y_test shape: {y_test.shape}")
```

```
X_train shape: (12793, 6)
X_test shape: (3199, 6)
y_train shape: (12793,)
y_test shape: (3199,)
```

```
[15]: import pandas as pd
      import numpy as np
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LinearRegression

      # Step 1: Load and preprocess the data
      # Assuming `data` is already loaded
```

```

# Replace 'Earnings_1978' with the desired target variable
X = data.drop(['Earnings_1974', 'Earnings_1975', 'Earnings_1978'], axis=1)
y = data['Earnings_1978']

# Step 2: Encode categorical variables
X = pd.get_dummies(X, drop_first=True) # Convert categorical columns to dummy_
    ↪variables

# Step 3: Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Step 4: Train the model
model = LinearRegression()
model.fit(X_train, y_train) # Fit the linear regression model

# Step 5: Evaluate the model
score = model.score(X_test, y_test) # R-squared score
print(f"R-squared score: {score}")

```

R-squared score: 0.09200104293622746

```

[17]: import joblib
y_pred = model.predict(X_test)
# Save the model
joblib.dump(model, 'linear_regression_model.pkl')

# Save predictions
predictions = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
predictions.to_csv('predictions.csv', index=False)

```